


Observation & reward pipeline

make_env(), src/agents/base.py — one agent step, RND-paper conditions



Conditions match Burda et al. 2018 (RND): sticky actions instead of random no-op starts, episode cap 4,500 agent steps (18,000 raw frames), full-episode training (no life-loss termination), reward clipped after true-score logging. RND networks see only the newest, whitened frame.